

Untitled5

December 18, 2023

1 Predicting Titanic Survival: A Data Science Project

This project involves developing a predictive model to determine the likelihood of survival for passengers on the Titanic using data science techniques in Python. The dataset encompasses various features such as socio-economic status, age, gender, and family size, providing valuable insights into factors influencing survival rates.

Dataset Columns:

PassengerId: Unique identifier for each passenger.

Survived: Binary indicator (1 or 0) for passenger survival.

Pclass: Ticket class representing socio-economic status (1st, 2nd, 3rd).

Name: Full name, including titles.

Sex: Gender of the passenger (male or female).

Age: Age of the passenger.

SibSp: Number of siblings/spouses aboard.

Parch: Number of parents/children aboard.

Ticket: Ticket number.

Fare: Passenger fare.

Cabin: Cabin number (with missing values).

Embarked: Port of embarkation (C, Q, S).

This project aims to leverage machine learning algorithms to analyze historical data and make predictions on passenger survival. By understanding the relationships between these features and survival outcomes, we can uncover patterns and insights that contribute to the overall narrative of the Titanic disaster.

1.1 Step 1: Import Libraries

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
```

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from xgboost import XGBClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix

```

1.2 Step 2: Load and Explore the Dataset

```

[3]: titanic_data = pd.read_csv('titanic.csv')
titanic_data.head()
titanic_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null   int64
1   Survived         891 non-null   int64
2   Pclass           891 non-null   int64
3   Name             891 non-null   object
4   Sex              891 non-null   object
5   Age              714 non-null   float64
6   SibSp            891 non-null   int64
7   Parch            891 non-null   int64
8   Ticket           891 non-null   object
9   Fare             891 non-null   float64
10  Cabin            204 non-null   object
11  Embarked         889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

```

[4]: titanic_data.describe()

```

```

[4]:
count    PassengerId    Survived    Pclass    Age    SibSp  \
count    891.000000    891.000000    891.000000    714.000000    891.000000
mean      446.000000      0.383838      2.308642     29.699118      0.523008
std       257.353842      0.486592      0.836071     14.526497      1.102743
min         1.000000      0.000000      1.000000      0.420000      0.000000
25%       223.500000      0.000000      2.000000     20.125000      0.000000
50%       446.000000      0.000000      3.000000     28.000000      0.000000
75%       668.500000      1.000000      3.000000     38.000000      1.000000
max       891.000000      1.000000      3.000000     80.000000      8.000000

```


| | Parch | Fare |
|-------|------------|------------|
| count | 891.000000 | 891.000000 |
| mean | 0.381594 | 32.204208 |
| std | 0.806057 | 49.693429 |
| min | 0.000000 | 0.000000 |
| 25% | 0.000000 | 7.910400 |
| 50% | 0.000000 | 14.454200 |
| 75% | 0.000000 | 31.000000 |
| max | 6.000000 | 512.329200 |

1.3 Step 3: Data Visualization

2 - Count Plot

```
[5]: sns.set(style="whitegrid")
plt.figure(figsize=(12, 6))
sns.countplot(x='Pclass', hue='Survived', data=titanic_data)
```


```
[5]: <Axes: xlabel='Pclass', ylabel='count'>
```



3 - Hist Plot


```
[6]: sns.histplot(titanic_data['Age'].dropna(), kde=False, bins=30)
```

```
[6]: <Axes: xlabel='Age', ylabel='Count'>
```



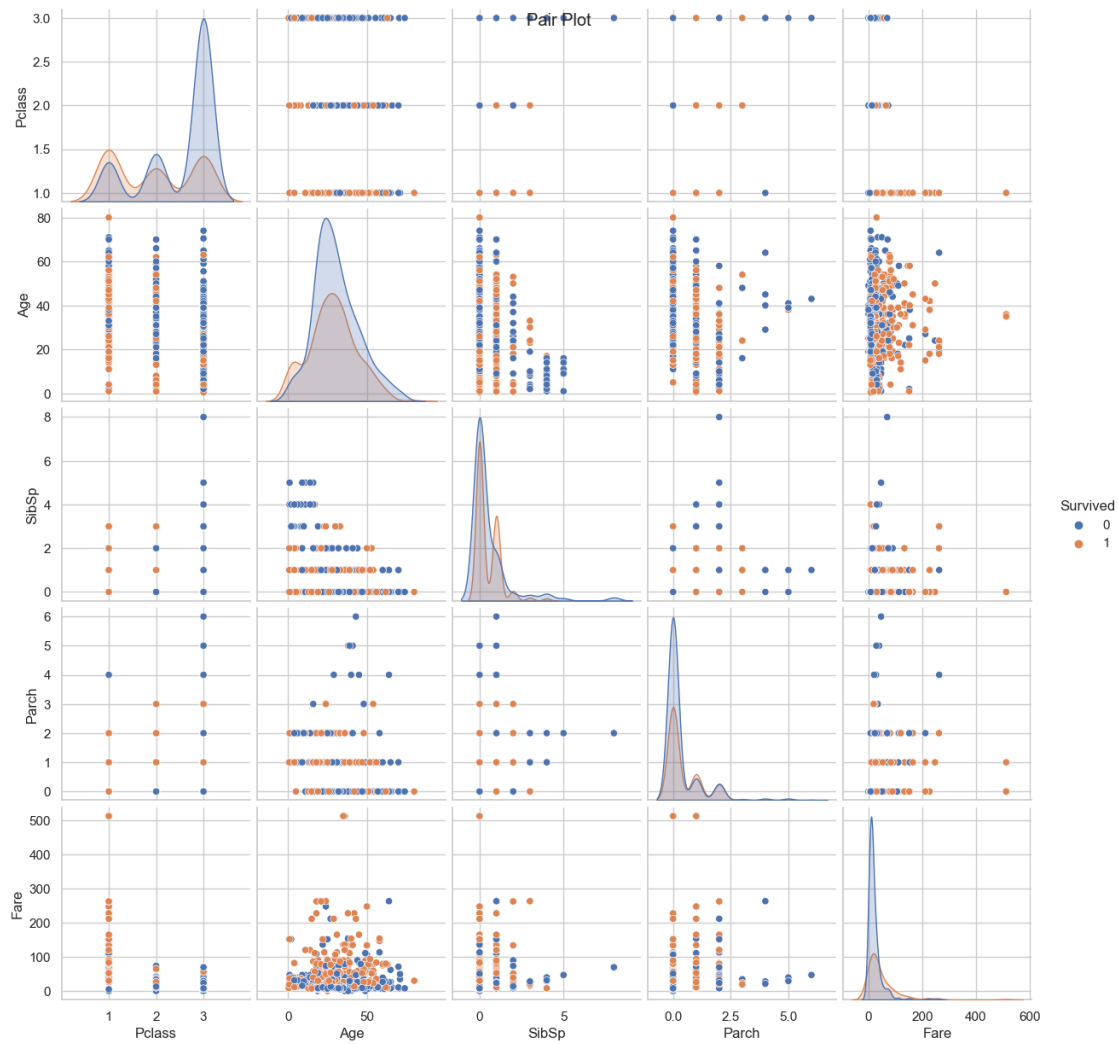
4 - Scatter Plot

```
[7]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Fare', hue='Survived', data=titanic_data, alpha=0.7)
plt.title('Scatter Plot: Age vs Fare')
plt.show()
```




5 - Pair Plot

```
[8]: sns.pairplot(titanic_data[['Survived', 'Pclass', 'Age', 'SibSp', 'Parch',  
    ↪ 'Fare']], hue='Survived')  
plt.suptitle('Pair Plot')  
plt.show()
```




6 - Dist Plot

```
[9]: plt.figure(figsize=(10, 6))
sns.histplot(titanic_data['Age'].dropna(), kde=True, bins=30, color='skyblue')
plt.title('Age Distribution')
plt.show()
```




7 - Bar Plot

```
[10]: plt.figure(figsize=(10, 6))
sns.countplot(x='Pclass', data=titanic_data, palette='viridis')
plt.title('Number of Passengers by Class')
plt.show()
```



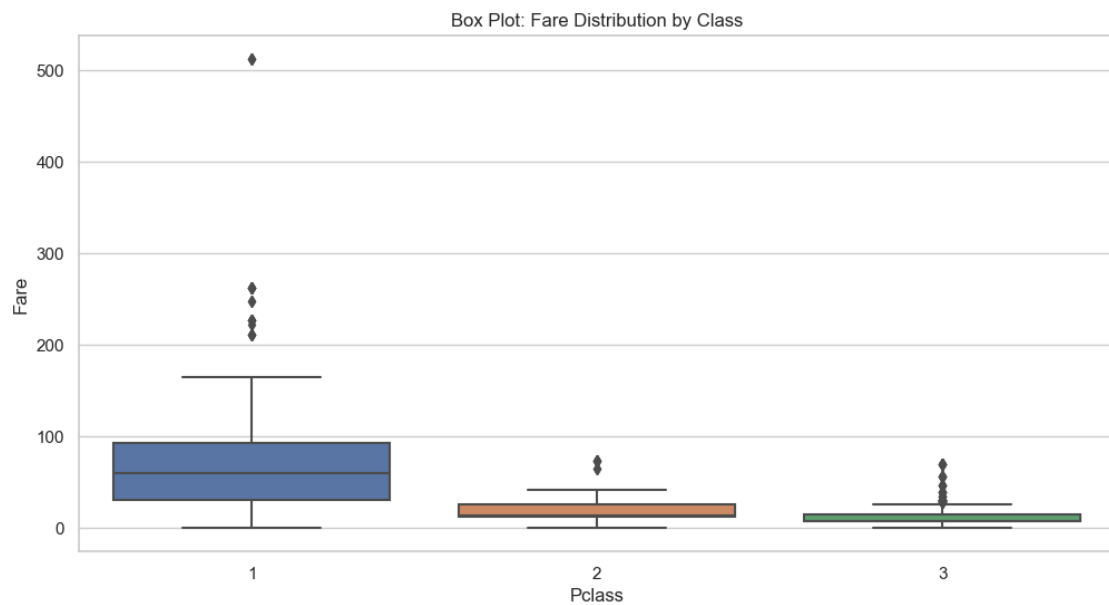
8 - Violin Plot

```
[11]: plt.figure(figsize=(12, 6))
sns.violinplot(x='Pclass', y='Age', hue='Survived', data=titanic_data,
              ↪split=True)
plt.title('Violin Plot: Age Distribution by Class')
plt.show()
```

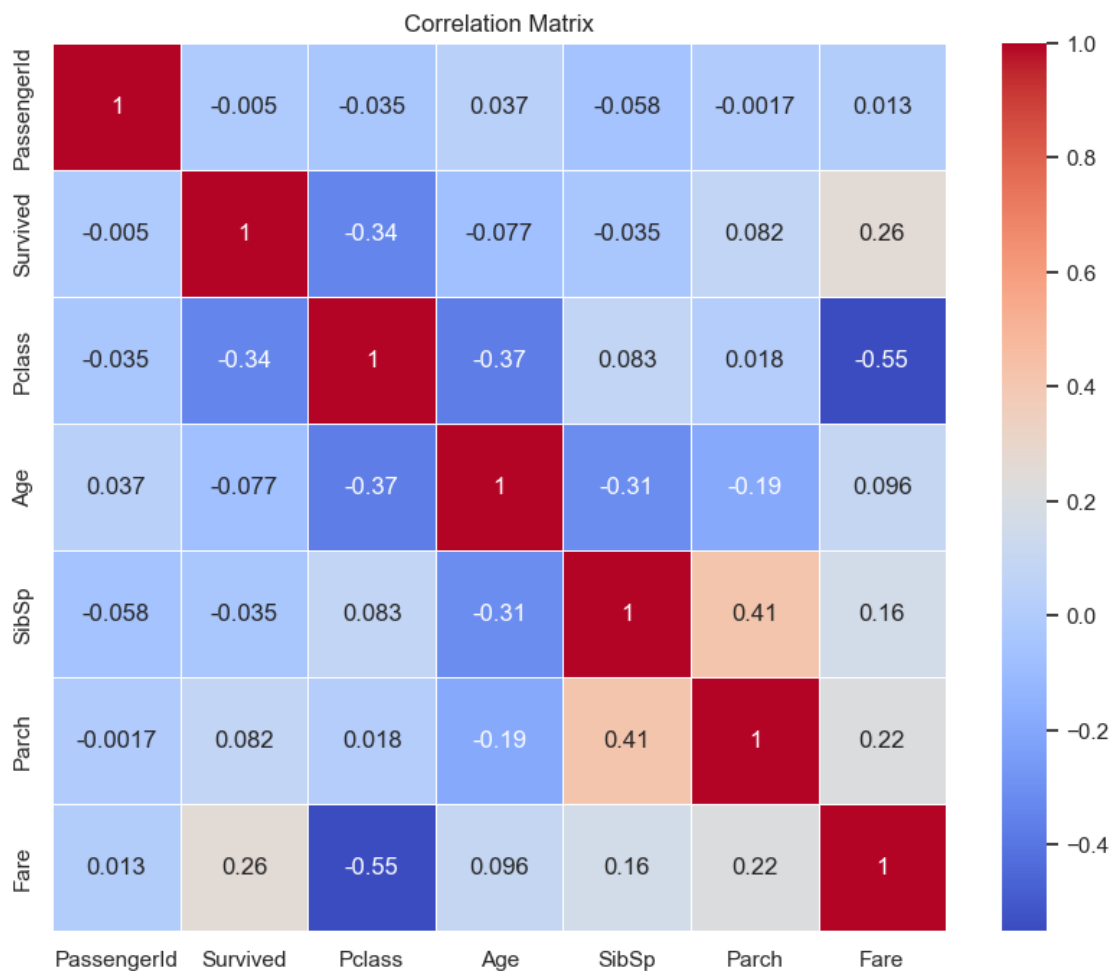
9 - Box Plot

```
[12]: plt.figure(figsize=(12, 6))
sns.boxplot(x='Pclass', y='Fare', data=titanic_data)
plt.title('Box Plot: Fare Distribution by Class')
plt.show()
```



10 - HeatMap

```
[13]: correlation_matrix = titanic_data.corr()
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=.5)
plt.title('Correlation Matrix')
plt.show()
```



10.1 Step 4: Data Preprocessing

```
[14]: titanic_data.drop(['Cabin'], axis=1, inplace=True)
titanic_data['Age'].fillna(titanic_data['Age'].median(), inplace=True)
titanic_data['Embarked'].fillna(titanic_data['Embarked'].mode()[0],
    ↪inplace=True)
```

```
titanic_data = pd.get_dummies(titanic_data, columns=['Sex', 'Embarked'],
    ↪drop_first=True)

X = titanic_data.drop(['Survived', 'PassengerId', 'Name', 'Ticket'], axis=1)
y = titanic_data['Survived']
```

10.2 Step 5: Train-Test Split

```
[15]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
```

10.3 Step 6: Build and Evaluate Models

```
[16]: def train_and_evaluate_model(model, X_train, y_train, X_test, y_test):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    return accuracy

# Logistic Regression
logreg_model = LogisticRegression()
logreg_accuracy = train_and_evaluate_model(logreg_model, X_train, y_train,
    ↪X_test, y_test)

# K Nearest Neighbour
knn_model = KNeighborsClassifier()
knn_accuracy = train_and_evaluate_model(knn_model, X_train, y_train, X_test,
    ↪y_test)

# Random Forest
rf_model = RandomForestClassifier()
rf_accuracy = train_and_evaluate_model(rf_model, X_train, y_train, X_test,
    ↪y_test)

# Xtreme Boosting
xgb_model = XGBClassifier()
xgb_accuracy = train_and_evaluate_model(xgb_model, X_train, y_train, X_test,
    ↪y_test)

# Support Vector Machine
svm_model = SVC()
svm_accuracy = train_and_evaluate_model(svm_model, X_train, y_train, X_test,
    ↪y_test)

# Display accuracy scores
print("Logistic Regression Accuracy:", logreg_accuracy)
```

```

print("K Nearest Neighbour Accuracy:", knn_accuracy)
print("Random Forest Accuracy:", rf_accuracy)
print("Xtreme Boosting Accuracy:", xgb_accuracy)
print("Support Vector Machine Accuracy:", svm_accuracy)

```

Logistic Regression Accuracy: 0.7988826815642458
 K Nearest Neighbour Accuracy: 0.7150837988826816
 Random Forest Accuracy: 0.8156424581005587
 Xtreme Boosting Accuracy: 0.8212290502793296
 Support Vector Machine Accuracy: 0.6536312849162011

```

[17]: best_model = max([(logreg_accuracy, 'Logistic Regression'),
                        (knn_accuracy, 'K Nearest Neighbour'),
                        (rf_accuracy, 'Random Forest'),
                        (xgb_accuracy, 'Xtreme Boosting'),
                        (svm_accuracy, 'Support Vector Machine')])

print(f"The best model is {best_model[1]} with an accuracy of {best_model[0]:.2f}")

```

The best model is Xtreme Boosting with an accuracy of 0.82

10.4 Step 7: Make Predictions

```

[18]: # Assuming you have chosen XGBClassifier as the best model
best_model_instance = XGBClassifier()
best_model_instance.fit(X_train, y_train)

# Make predictions on the test set
y_pred = best_model_instance.predict(X_test)

```

10.5 Step 8: Evaluate the Model

```

[19]: columns_for_prediction = ['PassengerId', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare', 'Sex_male', 'Embarked_Q', 'Embarked_S']
titanic_data_for_prediction = titanic_data[columns_for_prediction]

likelihood_of_survival = best_model_instance.predict(titanic_data_for_prediction.drop(['PassengerId'], axis=1))

print("Likelihood of Survival:", likelihood_of_survival)

```

Likelihood of Survival: [0 1 1 1 0 0 0 0 1 1 1 1 0 0 0 1 0 0 0 1 0 0 1 0 0 0 0
 1 0 0 1 1 0 0 0 0
 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 1 0 0 1 0 1 0 0 1 0 0 0 0 1 0 1 0 0 0 0 0
 1 0 0 0 1 1 0 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 1 0 0 0 0
 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 1 0 1 1 0 0 0 0 1 0 0 1 1 0 0 0 1 1 0 0 0 1 0
 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 1 0 1 0 0 0 0 0 0 0 0 1 1

```

0 1 0 0 0 1 0 0 1 1 1 0 0 1 1 0 0 0 0 0 0 0 1 1 0 0 1 0 0 0 1 1 0 1 0 0 0
0 0 1 0 0 0 0 0 1 1 0 1 0 1 0 1 0 0 0 0 1 0 0 0 0 0 1 1 0 0 0 0 0 0 1 1 1 1
1 0 1 0 0 1 0 0 0 1 1 0 0 1 0 1 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 0 0
0 0 1 1 1 1 0 1 0 1 1 1 0 1 1 1 1 0 0 1 1 0 1 1 0 0 1 1 0 1 0 1 1 1 1 0 0
0 1 0 0 1 0 0 1 1 0 0 0 1 1 1 1 0 0 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 1 1 1 1
1 0 0 0 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 1 1 0 0 1 1 0 1 0 0 1 0 0 0 0 0 0 0
1 0 0 0 0 1 0 1 0 1 1 0 0 0 0 0 0 0 0 0 1 1 0 0 0 1 1 0 0 1 0 1 0 0 1 0 0 1
0 1 1 1 1 1 0 0 0 1 0 1 0 1 1 0 1 0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 0 0 1 0
0 0 1 1 0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 1 0 0 1 1 0 1 1 0 0 1 0
1 0 1 0 0 1 0 0 1 0 0 0 1 0 0 1 0 1 0 1 0 1 1 0 0 1 0 0 1 1 0 1 1 0 0 1 1
0 1 0 1 1 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 0 1 1 0 1 1 1 0 0 0 1 0 1 0 0 0 1
0 0 0 0 1 0 0 1 1 0 0 0 1 0 0 1 1 1 0 0 1 0 0 1 0 0 1 0 0 1 1 0 0 0 0 1 0
0 1 0 1 0 0 1 1 0 0 0 0 1 0 0 1 1 0 1 0 1 0 1 0 1 0 0 0 1 0 1 1 0 0 0 1 0
0 0 0 1 1 0 0 0 0 0 0 1 0 1 0 1 0 0 0 0 0 0 0 1 1 1 0 0 0 0 0 1 0 0 1 1 0
0 0 0 1 1 1 0 1 0 0 0 0 0 1 1 0 0 1 0 0 0 1 0 1 1 0 0 1 0 0 0 0 0 0 1 0 0
1 0 1 0 1 0 0 1 0 0 1 1 0 0 1 1 0 0 0 1 0 0 1 1 0 1 0 1 0 0 0 0 1 0 1 0 0
1 0 1 1 1 0 0 0 0 1 0 1 0 0 0 0 1 0 0 1 1 0 0 0 1 1 1 1 0 0 0 0 1 0 0 0 0
0 0 0 0 0 0 1 1 0 1 0 0 0 1 0 1 0 1 0 0 0 1 0 0 1 1 0 0 1 0 0 0 0 0 0 1 0
0 0 1 1 1 1 1 1 0 0 0 1 0 0 1 1 0 0 1 0 1 0 0 1 1 0 0 0 1 1 0 0 0 0 0 0 1
0 1 0]

```